

# Weather Prediction Analytics using Snowflake, Airflow, and dbt

**Course:** SJSU DATA 226 **Team members:** Yunseo Oh (020746206), Isabella Shi (019256965)

**Repository:** <https://github.com/yunseoyunnioh/data226-weather-analytics> **Date:** September 29, 2026

---

## 1. Problem Statement

---

Weather affects planning decisions -- commuting, travel, agriculture, retail demand -- but raw daily weather readings (a max/min temperature and a precipitation total) don't answer questions like "is this warmer than usual" or "how long has it been dry." This project builds a pipeline that pulls daily weather for two cities, Seoul and Toronto, ingests it into a cloud data warehouse, derives multi-day analytics from it (moving averages, anomalies, rolling precipitation, dry spell length), and exposes those analytics through a BI dashboard -- so a viewer can compare the two cities' current weather trend at a glance instead of reading raw numbers.

A secondary goal, driven by the assignment itself, was architectural: build **one** pipeline that treats "which city" as configuration/data rather than duplicating the ETL and transformation logic per city.

## 2. Requirements

---

1. Ingest daily weather for two cities (Seoul, Toronto) from the Open-Meteo API ( <https://api.open-meteo.com/v1/forecast> ).
2. Load it into Snowflake with a transaction-safe, idempotent write path (safe to rerun without creating duplicate rows).
3. Orchestrate the ingestion with Apache Airflow, using Airflow Variables for city configuration and an Airflow Connection for Snowflake credentials -- no hardcoded secrets.
4. Transform the raw data with dbt into meaningful, per-city analytics (moving averages, anomaly, rolling precipitation, dry spell length), with schema tests and a snapshot.
5. Trigger the dbt run from Airflow, strictly after the ETL succeeds.
6. Visualize the resulting analytics in a BI tool (Superset, Preset, or Tableau) with at least two screenshots showing the dashboard responding to a filter change.
7. Ship one GitHub repository containing all code, with DAGs under `dags/` and the dbt project under `dbt/` , plus a README and this report.


### 3. Specifications

| Component     | Choice                                                                                                                            | Why                                                                                                                                                                                    |
|---------------|-----------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Source API    | Open-Meteo /v1/forecast , past_days=60 , forecast_days=1                                                                          | Required endpoint; past_days gives enough trailing history in one request per city for 7-day windows and an anomaly baseline, and forecast_days=1 includes today, for 61 days per city |
| City config   | Airflow Variable CITY_CONFIG (JSON: name -> lat/lon/timezone)                                                                     | Adding a city means adding one JSON entry, no code change                                                                                                                              |
| Credentials   | Airflow Connection snowflake_conn , shared by both DAGs                                                                           | One place for all Snowflake settings; the dbt DAG passes them to dbt as DBT_* environment variables, so no account, database, or secret is written in the repo                         |
| Orchestration | Two DAGs: weather_etl_dag (schedule) -> weather_dbt_dag (triggered, schedule=None)                                                | Makes the ETL-before-dbt dependency explicit and inspectable in the Airflow UI, not just implied by cron timing                                                                        |
| Raw table     | RAW.CITY_WEATHER_DAILY , PK (city, date)                                                                                          | New table name, distinct from the single-city HW2/HW3 WEATHER_DAILY table, to avoid touching already-graded data with an incompatible schema                                           |
| Idempotency   | Full refresh: DELETE + INSERT inside BEGIN / COMMIT , ROLLBACK + re-raise on failure                                              | Every run leaves the table holding exactly the current 61-day window per city, so reruns cannot create duplicates                                                                      |
| dbt layers    | weather_daily (ephemeral) -> weather_metrics (ephemeral) -> weather_analytics (table), in models/transform/ and models/analytics/ |                                                                                                                                                                                        |
| Snapshot      | weather_snapshot , strategy='check' on the raw weather columns of weather_analytics                                               | No natural updated_at column exists, and Open-Meteo revises very recent days as more observations arrive, so check is more appropriate than timestamp                                  |
| BI tool       | <a href="#">Preset</a> (cloud-hosted Apache Superset)                                                                             | Free tier, no local install/Docker needed, connects directly to Snowflake                                                                                                              |

### 4. Architecture

Full Mermaid source and narrative notes: docs/architecture.md . Rendered diagram:

**System Architecture Diagram.** Open-Meteo -> Airflow ETL DAG -> Snowflake RAW -> (triggered) Airflow dbt DAG -> dbt transform/analytics -> Snowflake analytics layer -> Preset dashboard, with the CITY\_CONFIG Variable feeding the ETL DAG, the snowflake\_conn Connection feeding both DAGs, and the weather\_snapshot snapshot sitting alongside the analytics table.



System architecture diagram: Open-Meteo API through Airflow ETL, Snowflake RAW, the triggered dbt DAG, dbt's transform and analytics layers, the Snowflake snapshot, and the Preset dashboard




Figure 1. weather\_etl\_dag graph view: extract, transform, load, trigger\_dbt\_dag all succeeded




Figure 2. weather\_dbt\_dag graph view: dbt\_run, dbt\_test, dbt\_snapshot all succeeded

## 5. Data Flow

1. `weather_etl_dag.extract()` reads `CITY_CONFIG`, calls Open-Meteo once per city, and returns each city's daily data.
2. `weather_etl_dag.transform()` flattens each city's daily arrays into `(city, latitude, longitude, date, temp_max, temp_min, precipitation, weather_code)` rows.
3. `weather_etl_dag.load()` replaces the contents of `RAW.CITY_WEATHER_DAILY` with those rows (`DELETE + INSERT`) inside one transaction, then `trigger_dbt_dag` starts `weather_dbt_dag`.
4. `weather_dbt_dag` runs `dbt run` (builds `weather_analytics`, with `weather_daily` and `weather_metrics`), then `dbt test`, then `dbt snapshot` (`weather_snapshot`).
5. The BI tool queries `ANALYTICS.WEATHER_ANALYTICS` directly.

## 6. Airflow Implementation

- **DAG 1 -- weather\_etl\_dag** (`dags/weather_etl_dag.py`): TaskFlow API (`@task`), tasks `extract -> transform -> load`, scheduled daily (`30 2 * * *`), `catchup=False`, `max_active_runs=1`. Final task `trigger_dbt_dag` (`TriggerDagRunOperator`) starts `weather_dbt_dag` after `load` succeeds.
- **DAG 2 -- weather\_dbt\_dag** (`dags/weather_dbt_dag.py`): three `BashOperator` tasks, `dbt_run -> dbt_test -> dbt_snapshot`, each invoking `dbt <command> --profiles-dir /opt/airflow/dbt` against the dbt project mounted into the container. `schedule=None`: this DAG never runs on its own timer, only via the trigger from DAG 1.
- **Airflow Variable**: `CITY_CONFIG`, a JSON object mapping city name to `{latitude, longitude, timezone}`.

- **Airflow Connection:** `snowflake_conn` , type Snowflake, key-pair authentication (private key file mounted read-only into the container; passphrase stored only in the Connection's Password field, never in code or files). The ETL DAG uses it through `SnowflakeHook` . The dbt DAG reads it with `BaseHook.get_connection('snowflake_conn')` and passes the values to dbt as `DBT_*` environment variables, which `dbt/profiles.yml` reads with `env_var()` .
- **Exception handling:** `load()` wraps the transaction in `try/except` ; on any exception it issues `ROLLBACK` , prints the error, and **re-raises** so the Airflow task is correctly marked failed (no silent swallowing). `extract()` calls `raise_for_status()` on each API response, so an HTTP error fails the task before anything is loaded.




Figure 3. Airflow Admin > Variables: CITY\_CONFIG with the Seoul/Toronto JSON




Figure 4. Airflow Admin > Connections: snowflake\_conn (type "snowflake"), with no host/port/secret values exposed

```

14  SELECT city,
15         COUNT(*) AS row_count,
16         MIN(date) AS first_date,
17         MAX(date) AS last_date,
18         MAX(loaded_at) AS last_loaded
19  FROM RAW.CITY_WEATHER_DAILY
20  GROUP BY city
21  ORDER BY city;

```

Results (1 minute ago)

|   | CITY    | ROW_COUNT | FIRST_DATE | LAST_DATE  | LAST_LOADED             |
|---|---------|-----------|------------|------------|-------------------------|
| 1 | Seoul   | 61        | 2026-07-31 | 2026-09-29 | 2026-09-28 19:30:58.166 |
| 2 | Toronto | 61        | 2026-07-30 | 2026-09-28 | 2026-09-28 19:31:52.348 |

Figure 5. RAW.CITY\_WEATHER\_DAILY row counts per city after the 2026-09-29 scheduled load: Seoul 61 rows (2026-07-31 to 2026-09-29), Toronto 61 rows (2026-07-30 to 2026-09-28). The full-refresh load always leaves exactly 61 rows per city

## 7. Snowflake Design

All objects live in the database set on the `snowflake_conn` Connection; no database name appears in the DAGs or the dbt project. Schema `RAW` holds the landing table, schema `ANALYTICS` holds dbt's output, and schema `SNAPSHOTS` holds the dbt snapshot. The ETL DAG creates `RAW` and the raw table if they don't exist, and dbt creates `ANALYTICS` and `SNAPSHOTS`, so the pipeline runs on a database with no existing schemas. See `sql/create_tables.sql` for the raw table DDL (also embedded in `weather_etl_dag.py`'s `load()` so the pipeline is self-provisioning).

## 8. Table Structures

All three tables below were inspected directly in the live Snowflake account ( `DESC TABLE` ) to confirm actual column types and nullability, not assumed from the DDL/model SQL alone -- Snowflake sometimes normalizes or relaxes both (e.g. `INTEGER` is stored as `NUMBER(38,0)` ; `CREATE TABLE AS SELECT` makes every output column nullable regardless of the source). The two transform models, `weather_daily` and `weather_metrics` , are ephemeral, so they have no table or view in Snowflake.

### `RAW.CITY_WEATHER_DAILY` (raw, loaded by Airflow)

| Column        | Type             | Nullable | Constraint                               | Description                                                                                                                   |
|---------------|------------------|----------|------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------|
| CITY          | VARCHAR(64)      | N        | part of PK                               | City name ( <code>Seoul</code> , <code>Toronto</code> )                                                                       |
| LATITUDE      | FLOAT            | N        | --                                       | Requested latitude                                                                                                            |
| LONGITUDE     | FLOAT            | N        | --                                       | Requested longitude                                                                                                           |
| DATE          | DATE             | N        | part of PK                               | Local calendar date of the reading                                                                                            |
| TEMP_MAX      | FLOAT            | Y        | --                                       | Daily max temperature, Celsius                                                                                                |
| TEMP_MIN      | FLOAT            | Y        | --                                       | Daily min temperature, Celsius                                                                                                |
| PRECIPITATION | FLOAT            | Y        | --                                       | Daily precipitation total, mm                                                                                                 |
| WEATHER_CODE  | NUMBER(38,0)     | Y        | --                                       | WMO weather interpretation code (declared <code>INTEGER</code> ; Snowflake stores all integers as <code>NUMBER(38,0)</code> ) |
| LOADED_AT     | TIMESTAMP_NTZ(9) | Y        | default <code>CURRENT_TIMESTAMP()</code> | Time the row was inserted by the most recent load                                                                             |

Declared key: `PRIMARY KEY (CITY, DATE)` (declarative only in Snowflake; the full-refresh load means each `(city, date)` is inserted exactly once per run, and the `unique` test on `city_date` in `weather_analytics` verifies it downstream).

## ANALYTICS.WEATHER\_ANALYTICS (dbt table -- BI source)

| Column             | Type              | Nullable | Description                                                                                                          |
|--------------------|-------------------|----------|----------------------------------------------------------------------------------------------------------------------|
| CITY               | VARCHAR(64)       | Y        | Partition key for every window function                                                                              |
| LATITUDE           | FLOAT             | Y        | Carried through                                                                                                      |
| LONGITUDE          | FLOAT             | Y        | Carried through                                                                                                      |
| WEATHER_DATE       | DATE              | Y        | Calendar date (renamed from raw's DATE in weather_daily)                                                             |
| CITY_DATE          | VARCHAR(16777216) | Y        | Surrogate key CITY \ \  '_' \ \  TO_VARCHAR( DATE, 'YYYY-MM-DD' ), built in weather_daily ; tested unique + not_null |
| TEMP_MAX           | FLOAT             | Y        | Carried through                                                                                                      |
| TEMP_MIN           | FLOAT             | Y        | Carried through                                                                                                      |
| PRECIPITATION      | FLOAT             | Y        | Carried through                                                                                                      |
| WEATHER_CODE       | NUMBER(38,0)      | Y        | Carried through                                                                                                      |
| DAILY_AVG_TEMP     | FLOAT             | Y        | ( TEMP_MAX + TEMP_MIN ) / 2                                                                                          |
| DRY_SPELL_LENGTH   | NUMBER(18,0)      | Y        | Consecutive dry days (< 0.1mm precip) ending on this date                                                            |
| MOVING_AVG_TEMP_7D | FLOAT             | Y        | 7-day trailing average of DAILY_AVG_TEMP , PARTITION BY city , rounded to 2 decimals                                 |
| TEMP_ANOMALY       | FLOAT             | Y        | DAILY_AVG_TEMP minus that city's average over the loaded window, rounded to 2 decimals                               |
| ROLLING_PRECIP_7D  | FLOAT             | Y        | 7-day trailing sum of PRECIPITATION , PARTITION BY city , rounded to 2 decimals                                      |

No declared key (table, materialized from a SELECT ). Every column is nullable at the database level -- CREATE TABLE AS SELECT does not carry NOT NULL forward from the raw table, even though the city\_date / city / weather\_date columns are never actually null in practice. Uniqueness and non-nullness are enforced by the dbt tests in models/schema.yml , not by the table's DDL.

## SNAPSHOTS.WEATHER\_SNAPSHOT (dbt snapshot)

| Column             | Type              | Nullable | Description                                                         |
|--------------------|-------------------|----------|---------------------------------------------------------------------|
| CITY               | VARCHAR(64)       | Y        | From weather_analytics                                              |
| LATITUDE           | FLOAT             | Y        | From weather_analytics                                              |
| LONGITUDE          | FLOAT             | Y        | From weather_analytics                                              |
| WEATHER_DATE       | DATE              | Y        | From weather_analytics                                              |
| CITY_DATE          | VARCHAR(16777216) | Y        | Snapshot unique_key                                                 |
| TEMP_MAX           | FLOAT             | Y        | Tracked by check strategy                                           |
| TEMP_MIN           | FLOAT             | Y        | Tracked by check strategy                                           |
| PRECIPITATION      | FLOAT             | Y        | Tracked by check strategy                                           |
| WEATHER_CODE       | NUMBER(38,0)      | Y        | Tracked by check strategy                                           |
| DAILY_AVG_TEMP     | FLOAT             | Y        | From weather_analytics , not a check_cols column                    |
| DRY_SPELL_LENGTH   | NUMBER(18,0)      | Y        | From weather_analytics , not a check_cols column                    |
| MOVING_AVG_TEMP_7D | FLOAT             | Y        | From weather_analytics , not a check_cols column                    |
| TEMP_ANOMALY       | FLOAT             | Y        | From weather_analytics , not a check_cols column                    |
| ROLLING_PRECIP_7D  | FLOAT             | Y        | From weather_analytics , not a check_cols column                    |
| DBT_SCD_ID         | VARCHAR(32)       | Y        | dbt-generated: unique ID for this SCD Type 2 row version            |
| DBT_UPDATED_AT     | TIMESTAMP_NTZ(9)  | Y        | dbt-generated: when this row version was recorded                   |
| DBT_VALID_FROM     | TIMESTAMP_NTZ(9)  | Y        | dbt-generated: start of this row version's validity                 |
| DBT_VALID_TO       | TIMESTAMP_NTZ(9)  | Y        | dbt-generated: end of this row version's validity ( NULL = current) |

Snapshot key: `unique_key='city_date'` , `strategy='check'` on `[temp_max, temp_min, precipitation, weather_code]` (see `snapshots/weather_snapshot.sql` ). The computed metric columns are stored in every row version but are not in `check_cols` ; Section 13 explains why.

## 9. Idempotency Design

Rerunning `weather_etl_dag` must not create duplicate rows. This is enforced in `load()` (`dags/weather_etl_dag.py`) with full-refresh pattern:

1. `CREATE SCHEMA IF NOT EXISTS RAW` and `CREATE TABLE IF NOT EXISTS` run **before** `BEGIN` , since Snowflake DDL auto-commits and would otherwise end an open transaction early.

2. Inside one `BEGIN`, `DELETE FROM RAW.CITY_WEATHER_DAILY` removes the previous load, and every extracted row is inserted with a parameterized `INSERT (%s` placeholders, so a missing API value becomes SQL `NULL`).
3. `COMMIT` makes the delete and inserts visible together. Any exception anywhere in the block triggers `ROLLBACK` and re-raises, so the previous data is kept and Airflow marks the task failed.

Net effect: running the DAG twice for the same date updates that date's row in place rather than appending a second copy.

## 10. dbt Implementation

---

See `dbt/README.md` for the full layer-by-layer breakdown. Summary: `weather_daily` (ephemeral) -> `weather_metrics` (ephemeral) -> `weather_analytics` (table), all analytics window functions `PARTITION BY city`.

## 11. dbt Models

---

- `models/transform/weather_daily.sql`
- `models/transform/weather_metrics.sql`
- `models/analytics/weather_analytics.sql`

## 12. dbt Tests

---

Defined in `models/schema.yml`, all on `weather_analytics`: `unique + not_null on city_date`; `not_null on city` and `weather_date`; `accepted_values on city` (Seoul, Toronto). Because the transform models pass every row straight through to `weather_analytics`, testing the final table also covers them.




Figure 6. weather\_dbt\_dag "dbt\_run" task log: weather\_analytics table built successfully, PASS=1 WARN=0 ERROR=0




Figure 7. weather\_dbt\_dag "dbt\_test" task log: all 5 tests individually PASS, Done. PASS=5 WARN=0 ERROR=0 SKIP=0 TOTAL=5

## 13. dbt Snapshot

snapshots/weather\_snapshot.sql , strategy='check' , keyed on city\_date , with check\_cols limited to the raw weather columns ( temp\_max , temp\_min , precipitation , weather\_code ). See Section 8 for its resulting columns and Section 3 for why check was chosen over timestamp .

The computed metrics are left out of `check_cols` on purpose. They shift whenever the 61-day window moves even if the weather data hasn't changed: `temp_anomaly`, for example, is measured against the city's average over the whole window, so when one day drops out and a new one comes in, every row's anomaly changes slightly. Checking those columns would create a new version of all 122 rows on every run and bury the real revisions. With only the raw columns checked, a new version is recorded when Open-Meteo revises a day's values, and a row is closed out when its day ages out of the window.

Both cases showed up in testing. After a few runs on 2026-09-26/27, the snapshot held 124 rows: 122 current rows, one earlier version of `Seoul_2026-09-27` (its `temp_max` was revised from 26.2 to 26.3 and `temp_min` from 16.3 to 16.2 as the day's observations came in), and one closed-out `Toronto_2026-07-28`, which dropped out of the window when Toronto's local date rolled over to September 27.




Figure 8. weather\_dbt\_dag "dbt\_snapshot" task log: 1 of 1 OK snapshot snapshots.weather\_snapshot

## 14. Airflow/dbt Scheduling

`weather_etl_dag` is the only DAG on a cron schedule (`30 2 * * *`). `weather_dbt_dag` has `schedule=None` and is started exclusively by `weather_etl_dag`'s last task, `trigger_dbt_dag` (`TriggerDagRunOperator(trigger_dag_id="weather_dbt_dag")`). This makes the "ETL succeeds, then dbt runs" dependency visible directly in the DAG code and in the Airflow UI (two separate DAGs, one triggering the other), rather than relying on scheduling two DAGs far enough apart in time.

## 15. BI Dashboard

---

**Status:** complete. **Tool:** [Preset](#) (cloud-hosted Apache Superset). **Team/Workspace:** "DATA 226 Weather Analytics" / "Weather Analytics Lab". **Dashboard name:** "Weather Analytics Dashboard". **Dataset:** ANALYTICS.WEATHER\_ANALYTICS .

**Purpose:** let a viewer compare Seoul vs. Toronto's current weather trend -- temperature level and trajectory, anomaly, and precipitation -- without reading raw numbers. **Usage:** select a city (or leave both selected) via the dashboard's CITY filter, narrow the WEATHER\_DATE time range, and all four charts update together.

**Authentication:** Preset could not use a plain Snowflake username/password connection because this account requires MFA/TOTP. The dashboard instead connects using the same Snowflake **key-pair authentication** already used by Airflow and dbt, configured through Preset's **Advanced -> Security -> Secure Extra** field rather than the basic login form. That field lives only in Preset's hosted connection config -- the key and passphrase were never written to this repository. See the root README.md -> "Security / secrets".

**The four charts, all grouped by CITY :**

1. **Seoul vs. Toronto Daily Avg Temp** -- line chart, X = WEATHER\_DATE , metric AVG(DAILY\_AVG\_TEMP) , dimension CITY .
2. **7-Day Moving Average Temperature** -- line chart, X = WEATHER\_DATE , metric AVG(MOVING\_AVG\_TEMP\_7D) , dimension CITY .
3. **Temperature Anomaly** -- bar chart, X = WEATHER\_DATE , metric AVG(TEMP\_ANOMALY) , dimension CITY .
4. **7-Day Rolling Precipitation** -- bar chart, X = WEATHER\_DATE , metric AVG(ROLLING\_PRECIP\_7D) , dimension CITY .

**Filters:** native dashboard filters on CITY and on WEATHER\_DATE (time range), applied across all four charts at once.




Figure 9. Weather Analytics Dashboard, default state (no filters applied): all four charts show both Seoul and Toronto over the full loaded window




Figure 10. Weather Analytics Dashboard filtered to City = Seoul and Date Range 2026-09-01 to 2026-09-20: all four charts update to a single city and a narrower window

## 16. Results / Observations

The pipeline was run end-to-end against live data (see Section 18, Conclusion, for the full verification summary). As of the run on 2026-09-25/26, over the ~61-day loaded window:

- Both cities were in the middle of an extended dry spell at the same time: Seoul had gone 14 consecutive days with under 0.1mm of precipitation, Toronto 17 days -- both `rolling_precip_7d` values were at or effectively at 0.
- Both cities' most recent days ran cooler than their loaded-window average: Seoul's temperature anomaly was around -3 to -5 degrees C over its last several days; Toronto's was around -5 to -8 degrees C, reaching -7.76 on 2026-09-22 (its `daily_avg_temp` that day, 12.15 degrees C, was well under its ~19.9 degrees C window average).
- Seoul's window average `daily_avg_temp` (~24.8 degrees C) was meaningfully warmer than Toronto's (~19.9 degrees C) over the same calendar period, as expected in late-summer-to-fall for the two cities' respective climates.
- Anomaly range across the window: Seoul -6.36 to +5.94 degrees C; Toronto -7.76 to +5.74 degrees C -- both cities show comparable volatility around their own baseline despite the ~5 degree gap in absolute average temperature, which is exactly the kind of city-relative comparison `temp_anomaly` (computed `PARTITION BY city`) is meant to surface.

These figures came from querying `ANALYTICS.WEATHER_ANALYTICS` directly after a real `dbt run`, and the Preset dashboard confirms the same pattern visually (Figures 9-10, Section 15): the default-view "Temperature Anomaly" chart ranges from roughly -8 to +6 for both cities, matching the -7.76 to +5.94 range computed here, and "7-Day Rolling Precipitation" shows both cities' bars taper toward zero at the right edge of the window, matching the Seoul/Toronto dry spell described above. The filtered view (Figure 10, isolated to Seoul, 2026-09-01 to 2026-09-20) shows Seoul's rolling precipitation dropping from roughly 150mm to near 0 by the end of that range and its anomaly staying mostly negative through September -- the same transition into the dry spell visible in the unfiltered chart, now isolated to one city.

## 17. Future Work

---

- Add more cities by extending the `CITY_CONFIG` Airflow Variable and the `accepted_values` test.
- Backfill a longer history (Open-Meteo's archive API) to make `temp_anomaly` a true climatological anomaly instead of a within-loaded-window baseline.
- Add a `freshness` check on the `raw.city_weather_daily` source so a stale load surfaces before `dbt` silently reruns on old data.
- Alerting (Slack/email) on Airflow task failure or `dbt` test failure.

## 18. Conclusion

---

The pipeline was verified end-to-end in a real local Airflow environment: `weather_etl_dag` (extract, transform, load) and the triggered `weather_dbt_dag` (`dbt` run, test, snapshot) both completed successfully against the live Snowflake account, loading real Seoul and Toronto data and passing all 5 `dbt` tests. Rerunning the ETL DAG confirmed the full-refresh load is idempotent: the raw and analytics tables held 122 rows (61 per city) after every run. The repository was also tested from a fresh clone with a new Airflow container against an empty Snowflake database: following only the README setup ( `.env`, the `CITY_CONFIG` Variable, and the `snowflake_conn` Connection), both DAGs ran green and created all

three tables with no file edits. One Airflow codebase and one dbt project handle both cities through the `CITY_CONFIG` Variable and `PARTITION BY city` window functions, satisfying the assignment's requirement to treat city as configuration rather than forking the pipeline per city.

The BI layer is also complete: a Preset dashboard ("Weather Analytics Dashboard") connects to `ANALYTICS.WEATHER_ANALYTICS` -- via Snowflake key-pair authentication through Preset's Secure Extra configuration, working around this account's MFA requirement -- and presents four charts (daily average temperature, 7-day moving average, temperature anomaly, 7-day rolling precipitation), all grouped by city and filterable by city and date range. End to end, the project satisfies Open-Meteo -> Airflow -> Snowflake RAW -> dbt -> Snowflake analytics -> BI dashboard with one shared pipeline for both cities, and all 10 required screenshots (Airflow Variables and Connections, both DAG graph views, the raw table row counts, the three dbt task logs, and both dashboard views) are captured and present in `screenshots/`, embedded throughout this report as Figures 1-10 (Sections 4, 6, 12, 13, and 15), alongside the system architecture diagram in Section 4. The repository is public at the URL above, and the submission date is recorded above. This report is complete.